Appendix A: Tutorial

A typical work-flow in the Python part of the framework is as follows:

- Load json with control flow graph into a `Module` object

- Perform necessary data-flow analysis

- Execute the chosen register allocation algorithm(s)

- Test and compare results

Below, we show an example of interactive, step-by-step register allocation based on the basic version of the Linear Scan algorithm. At the end, we will present the usage of full, end-to-end allocation procedure and compare this algorithm with the Extended Linear Scan. As an input, we will use a simple program computing the greatest common divisor.

1 Loading programs from json and data flow analysis.

We assume to already have a json file with the Control Flow Graph of our program, generated by the LLVM plugin. It can be now loaded into a `Module` object which will give us access to all the functions in the file. Here, we only have two functions: `gcd` and `main`. We will focus on the first.

Having loaded the file, we perform full data flow analysis on our `Module`, including liveness, dominance and loop analysis which is everything we need for register allocation. At the end, we print out the function to the output.

```
In [1]: import cfg
        import cfg.analysis as analysis
        from cfg.printer import FunctionString

        m = cfg.Module.from_file("programs/gcd.json")
        analysis.perform_full_analysis(m)

        print "Functions in the module: ", ", ".join(m.functions.keys()), "\n"

        gcd = m.functions['gcd']
        print FunctionString(gcd)

Functions in the module:  main, gcd

bb1(entry)
  0:   v1 = icmp v2 v3
  1:   v4 = br v1 bb3 bb2


bb2(if.then)
  2:   v5 = xor v2 v3
  3:   v6 = xor v3 v5
  4:   v7 = xor v5 v6
  5:   v8 = br bb3


bb3(if.end)
  6:   v9 = phi bb2 -> v6 bb1 -> v3
  7:   v10 = phi bb2 -> v7 bb1 -> v2
```

```
  8:  v11 = br bb4
```

```
bb4(while.cond)
  9:  v12 = phi bb5 -> v13 bb3 -> v9
 10:  v14 = phi bb5 -> v12 bb3 -> v10
 11:  v15 = icmp v12 const
 12:  v16 = br v15 bb6 bb5
```

```
bb5(while.body)
 13:  v13 = srem v14 v12
 14:  v17 = br bb4
```

```
bb6(while.end)
 15:  v18 = ret v14
```

The `cfg.printer` module contains helper classes for printing various objects we operate on, such as `FunctionString`, `BBString` or `InstructionString`. While using any of them, we can pass options saying what we want to include in the object description. Here, apart from the sole function body, we can print predecessors, successors, liveness sets, dominance sets and so on.

```
In [2]: from cfg.printer import Opts
        print FunctionString(gcd, Opts(predecessors=True, successors=True,
                                       liveness=True, dominance=True))
```

```
bb1(entry)
  0:   v1 = icmp v2 v3
  1:   v4 = br v1 bb3 bb2
     PREDS: []
     SUCCS: [bb3, bb2]
   LIVE-IN: [v2, v3]
  LIVE-OUT: [v2, v3]
       DOM: [bb1]
```

```
bb2(if.then)
  2:   v5 = xor v2 v3
  3:   v6 = xor v3 v5
  4:   v7 = xor v5 v6
  5:   v8 = br bb3
     PREDS: [bb1]
     SUCCS: [bb3]
   LIVE-IN: [v2, v3]
  LIVE-OUT: [v6, v7]
       DOM: [bb1, bb2]
```

```
bb3(if.end)
  6:   v9 = phi bb2 -> v6 bb1 -> v3
  7:  v10 = phi bb2 -> v7 bb1 -> v2
  8:  v11 = br bb4
     PREDS: [bb2, bb1]
     SUCCS: [bb4]
```

```
    LIVE-IN: [v9, v10]
   LIVE-OUT: [v9, v10]
        DOM: [bb1, bb3]


bb4(while.cond)
  9:  v12 = phi bb5 -> v13 bb3 -> v9
 10:  v14 = phi bb5 -> v12 bb3 -> v10
 11:  v15 = icmp v12 const
 12:  v16 = br v15 bb6 bb5
      PREDS: [bb5, bb3]
      SUCCS: [bb6, bb5]
    LIVE-IN: [v12, v14]
   LIVE-OUT: [v12, v14]
        DOM: [bb1, bb3, bb4]


bb5(while.body)
 13:  v13 = srem v14 v12
 14:  v17 = br bb4
      PREDS: [bb4]
      SUCCS: [bb4]
    LIVE-IN: [v12, v14]
   LIVE-OUT: [v12, v13]
        DOM: [bb1, bb3, bb4, bb5]


bb6(while.end)
 15:  v18 = ret v14
      PREDS: [bb4]
      SUCCS: []
    LIVE-IN: [v14]
   LIVE-OUT: []
        DOM: [bb1, bb3, bb4, bb6]
```

2 Linear Scan register allocation

2.1 Intervals building

The first stage of the Linear Scan algorithm is bulding lifetime intervals. Below, we create an object of
BasicLinearScan, build lifetime intervals from our function and print it out. We do it on a copy of the
function as the allocation may change its structure.

```
In [3]: g = gcd.copy()

        from allocators.lscan.basic import BasicLinearScan
        from cfg.printer import IntervalsString

        bls = BasicLinearScan()
        intervals = bls.compute_intervals(g)
        print IntervalsString(intervals)

INTERVAL              VAR-ID        REG
[-0.1, 2]             v2             -
```

```
[-0.1, 3]           v3              -
[0, 1]              v1              -
[2, 4]              v5              -
[3, 5.1]            v6              -
[4, 5.1]            v7              -
[5.9, 8.1]          v10             -
[5.9, 8.1]          v9              -
[8.9, 14.1]         v12             -
[8.9, 15]           v14             -
[11, 12]            v15             -
[13, 14.1]          v13             -
```

We can also draw the intervals on a chart in a following way:

```
In [4]: import utils
        %matplotlib inline
        utils.draw_intervals(intervals, figsize=(8, 8))
```


All the intervals are black because they do not have any regisers assigned yet.

2.2 Register allocation

Now, let us try to perform register allocation with 2 available registers. The function responsible for allocation takes as input intervals, number of registers and a boolean variable denoting whether we allow spilling or not. The function returns `True` if it succeeded to allocate registers to all variables without spilling, and `False`, otherwise.

```
In [5]: success = bls.allocate_registers(intervals, 2, spilling=True)
        print "Allocation succeeded without spilling: ", success, "\n"
        print IntervalsString(intervals)
        utils.draw_intervals(intervals, regcount=2, figsize=(8, 8))

Allocation succeeded without spilling:  False
```

INTERVAL	VAR-ID	REG
[-0.1, 2]	v2	reg2
[-0.1, 3]	v3	-
[0, 1]	v1	reg1
[2, 4]	v5	reg1
[3, 5.1]	v6	reg2
[4, 5.1]	v7	reg1
[5.9, 8.1]	v10	reg2
[5.9, 8.1]	v9	reg1
[8.9, 14.1]	v12	reg2
[8.9, 15]	v14	-
[11, 12]	v15	reg1
[13, 14.1]	v13	reg1

Lifetime intervals

As we can see, 2 registers are not enough - two variables had to be spilled into memory. In the picture above, intervals with the same color have the same registers assigned. Spilled intervals are distinguished by a dashed line. In the next step, for each spilled interval (variable), we have to insert `store` and `load` instructions in appropriate fragments of the code.

2.3 Spill code insertion

The procedure responsible for inserting spill code is independent from allocation algorithms. It takes the function as the only argument and modifies it accordingly. We can print the results afterwards.

```
In [6]: cfg.resolve.insert_spill_code(g)
        print FunctionString(g, Opts(mark_spill=True))

bb1(entry)
  0:  v23 = load_ mem(v3)
  1:   v1 = icmp v2 v23
  2:   v4 = br v1 bb3 bb2


bb2(if.then)
  3:  v21 = load_ mem(v3)
  4:   v5 = xor v2 v21
  5:  v22 = load_ mem(v3)
  6:   v6 = xor v22 v5
  7:   v7 = xor v5 v6
  8:   v8 = br bb3


bb3(if.end)
  9:   v9 = phi bb2 -> v6 bb1 -> v3
 10:  v10 = phi bb2 -> v7 bb1 -> v2
 11:  v11 = br bb4


bb4(while.cond)
 12:  v12 = phi bb5 -> v13 bb3 -> v9
 13:  v14 = phi bb5 -> v12 bb3 -> v10
 14:  v15 = icmp v12 const
 15:  v16 = br v15 bb6 bb5


bb5(while.body)
 16:  v20 = load_ mem(v14)
 17:  v13 = srem v20 v12
 18:  v17 = br bb4


bb6(while.end)
 19:  v19 = load_ mem(v14)
 20:  v18 = ret v19
```

The inserted `load` and `store` instructions are colored in violet. We also add an underscore symbol to distinguish them from the same instructions present in the original code. Inserting new instructions changes the structure of the CFG, so the data-flow analysis should be repeated. We follow the convention that each procedure modifying the structure of the function is responsible for performing the necessary analysis afterwords. Therefore, we do not have to do it on our own.

2.4 Repeated register allocation

After spilling two variables, the register pressure decreased, and we can try to repeat the allocation (but we do not allow spilling any more). We build the intervals again and execute the allocation algorithm.

```
In [7]: intervals = bls.compute_intervals(g)
        print IntervalsString(intervals)
        utils.draw_intervals(intervals, regcount=2, figsize=(8, 8))
```

INTERVAL	VAR-ID	REG
[-0.1, 4]	v2	-
[0, 1]	v23	-
[1, 2]	v1	-
[3, 4]	v21	-
[4, 7]	v5	-
[5, 6]	v22	-
[6, 8.1]	v6	-
[7, 8.1]	v7	-
[8.9, 11.1]	v10	-
[8.9, 11.1]	v9	-
[11.9, 18.1]	v12	-
[14, 15]	v15	-
[16, 17]	v20	-
[17, 18.1]	v13	-
[19, 20]	v19	-

8

```
In [8]: success = bls.allocate_registers(intervals, 2, spilling=False)
        print "Allocation succeeded without spilling: ", success, "\n"
        print IntervalsString(intervals)
        utils.draw_intervals(intervals, regcount=2, figsize=(8, 8))
```

```
Allocation succeeded without spilling:  True

INTERVAL                VAR-ID          REG
[-0.1, 4]               v2              reg2
[0, 1]                  v23             reg1
[1, 2]                  v1              reg1
[3, 4]                  v21             reg1
[4, 7]                  v5              reg2
[5, 6]                  v22             reg1
[6, 8.1]                v6              reg1
[7, 8.1]                v7              reg2
[8.9, 11.1]             v10             reg1
[8.9, 11.1]             v9              reg2
[11.9, 18.1]            v12             reg1
[14, 15]                v15             reg2
[16, 17]                v20             reg2
[17, 18.1]              v13             reg2
[19, 20]                v19             reg1
```


Lifetime intervals

Now the allocator finished successfully. Let us print out the function with the allocated registers next to the corresponding variables.

In [9]: print FunctionString(g, Opts(with_alloc=True, liveness=True))

```
bb1(entry)
   0: v23(reg1) = load_ mem(v3)
   1: v1(reg1) = icmp v2(reg2) v23(reg1)
   2:   v4 = br v1(reg1) bb3 bb2
    LIVE-IN: [(v2, reg2)]
    LIVE-OUT: [(v2, reg2)]


bb2(if.then)
   3: v21(reg1) = load_ mem(v3)
   4: v5(reg2) = xor v2(reg2) v21(reg1)
   5: v22(reg1) = load_ mem(v3)
   6: v6(reg1) = xor v22(reg1) v5(reg2)
   7: v7(reg2) = xor v5(reg2) v6(reg1)
   8:   v8 = br bb3
    LIVE-IN: [(v2, reg2)]
    LIVE-OUT: [(v6, reg1), (v7, reg2)]


bb3(if.end)
   9: v9(reg2) = phi bb2 -> v6(reg1) bb1 -> v3(mem(v3))
  10: v10(reg1) = phi bb2 -> v7(reg2) bb1 -> v2(reg2)
  11:   v11 = br bb4
    LIVE-IN: [(v9, reg2), (v10, reg1)]
    LIVE-OUT: [(v9, reg2), (v10, reg1)]


bb4(while.cond)
  12: v12(reg1) = phi bb5 -> v13(reg2) bb3 -> v9(reg2)
  13: v14(mem(v14)) = phi bb5 -> v12(reg1) bb3 -> v10(reg1)
  14: v15(reg2) = icmp v12(reg1) const
  15:   v16 = br v15(reg2) bb6 bb5
    LIVE-IN: [(v12, reg1)]
    LIVE-OUT: [(v12, reg1)]


bb5(while.body)
  16: v20(reg2) = load_ mem(v14)
  17: v13(reg2) = srem v20(reg2) v12(reg1)
  18:   v17 = br bb4
    LIVE-IN: [(v12, reg1)]
    LIVE-OUT: [(v12, reg1), (v13, reg2)]


bb6(while.end)
  19: v19(reg1) = load_ mem(v14)
  20:   v18 = ret v19(reg1)
```

```
   LIVE-IN: []
  LIVE-OUT: []
```

2.5   Translating out of SSA form

The last thing we have to do is to translate the program out of SSA form. Alike the spill code insertion, ϕ-elimination is also independent from the allocation algorithm. However, if there are memory-to-memory copies or mov-cycles, it can create additional variables which need a regsiter. Therefore, apart from the function instance, we also pass in the argument the total number of available registers. It returns True on success and False otherwise.

```
In [10]: success = cfg.resolve.eliminate_phi(g, 2)
         print "Phi elimination succeeded: ", success, "\n"
         print FunctionString(g, Opts(with_alloc=True))

Phi elimination succeeded:  True

bb1(entry)
  0: v23(reg1) = load_ mem(v3)
  1: v1(reg1) = icmp v2(reg2) v23(reg1)
  2: v4 = br v1(reg1) bb3 bb2


bb2(if.then)
  3: v21(reg1) = load_ mem(v3)
  4: v5(reg2) = xor v2(reg2) v21(reg1)
  5: v22(reg1) = load_ mem(v3)
  6: v6(reg1) = xor v22(reg1) v5(reg2)
  7: v7(reg2) = xor v5(reg2) v6(reg1)
  8: store_ mem(v24) v6(reg1)
  9: v10(reg1) = mov v7(reg2)
 10: v9(reg2) = load_ mem(v24)
 11: v8 = br bb3


bb7(None)
 12: v10(reg1) = mov v2(reg2)
 13: v9(reg2) = load_ mem(v3)
 14: br bb3


bb3(if.end)
 15: store_ mem(v14) v10(reg1)
 16: v12(reg1) = mov v9(reg2)
 17: v11 = br bb4


bb4(while.cond)
 18: v15(reg2) = icmp v12(reg1) const
 19: v16 = br v15(reg2) bb6 bb5


bb5(while.body)
 20: v20(reg2) = load_ mem(v14)
 21: v13(reg2) = srem v20(reg2) v12(reg1)
```

```
22: store_ mem(v14) v12(reg1)
23: v12(reg1) = mov v13(reg2)
24: v17 = br bb4


bb6(while.end)
25: v19(reg1) = load_ mem(v14)
26:  v18 = ret v19(reg1)
```

On the listing above we can see all the moves generated by ϕ-elimination, even those between variables sharing the same register. We can also notice a new basic block - *bb7*, created on the edge between blocks *bb1* and *bb3*, which had to be inserted because of *bb1* having multiple successors. When we print the function with `alloc_only=True` option, registers are showed instead of variables, and redundant instructions or basic blocks are skipped. By setting `mark_non_ssa=True` we will see instructions produced by the ϕ-elimination phase colored in violet. The variables without registers are those that are not used anywhere in the program (they are just defined), and therefore are skipped by the allocation.

```
In [11]: print FunctionString(g, Opts(alloc_only=True, mark_non_ssa=True))

bb1(entry)
  0: reg1 = load_ mem(v3)
  1: reg1 = icmp reg2 reg1
  2:   v4 = br reg1 bb3 bb2


bb2(if.then)
  3: reg1 = load_ mem(v3)
  4: reg2 = xor reg2 reg1
  5: reg1 = load_ mem(v3)
  6: reg1 = xor reg1 reg2
  7: reg2 = xor reg2 reg1
  8: store_ mem(v24) reg1
  9: reg1 = mov reg2
 10: reg2 = load_ mem(v24)
 11:   v8 = br bb3


bb7(None)
 12: reg1 = mov reg2
 13: reg2 = load_ mem(v3)
 14: br bb3


bb3(if.end)
 15: store_ mem(v14) reg1
 16: reg1 = mov reg2
 17:  v11 = br bb4


bb4(while.cond)
 18: reg2 = icmp reg1 const
 19:  v16 = br reg2 bb6 bb5


bb5(while.body)
```

```
20: reg2 = load_ mem(v14)
21: reg2 = srem reg2 reg1
22: store_ mem(v14) reg1
23: reg1 = mov reg2
24:  v17 = br bb4


bb6(while.end)
25: reg1 = load_ mem(v14)
26:  v18 = ret reg1
```

3   Correctness checks

After successful allocation, we may want to check if the allocation, as well as data-flow, are correct. We do it by using `cfg.sanity` module:

```
In [12]: import cfg.sanity as sanity
         print "Allocation is correct: ", sanity.allocation_is_correct(g)
         print "Data flow is correct: ", sanity.data_flow_is_correct(g, gcd)

Allocation is correct:  True
Data flow is correct:  True
```

4   Cost of allocation

At the end, we can calculate the cost of the register allocation. Our main cost calculator depends on 3 parameters: S, N and L meaning respectively: the cost of spill instructions, the cost of normal instructions and the loop penalty. The final cost of the allocation is a difference between the modified and the original function costs. Here, we can also print out the detailed view of the cost computation by using `cfg.printer.CostString` .

```
In [13]: from cost import MainCostCalculator
         from cfg.printer import CostString

         mcc = MainCostCalculator()
         print CostString(g, mcc), "\n"

         print "Final cost of the allocation: ", mcc.function_diff(g, gcd)

Main cost (S=2, N=1, L=10)
  LOOP      COST              INSTR
    0       2.0        0: reg1 = load_ mem(v3)
    0       1.0        1: reg1 = icmp reg2 reg1
    0       1.0        2:   v4 = br reg1 bb3 bb2
    0       2.0        3: reg1 = load_ mem(v3)
    0       1.0        4: reg2 = xor reg2 reg1
    0       2.0        5: reg1 = load_ mem(v3)
    0       1.0        6: reg1 = xor reg1 reg2
    0       1.0        7: reg2 = xor reg2 reg1
    0       2.0        8: store_ mem(v24) reg1
    0       1.0        9: reg1 = mov reg2
    0       2.0       10: reg2 = load_ mem(v24)
    0       1.0       11:   v8 = br bb3
    0       1.0       12: reg1 = mov reg2
```

```
     0        2.0      13: reg2 = load_ mem(v3)
     0        1.0      14: br bb3
     0        2.0      15: store_ mem(v14) reg1
     0        1.0      16: reg1 = mov reg2
     0        1.0      17:  v11 = br bb4
     1       10.0      18: reg2 = icmp reg1 const
     1       10.0      19:  v16 = br reg2 bb6 bb5
     1       20.0      20: reg2 = load_ mem(v14)
     1       10.0      21: reg2 = srem reg2 reg1
     1       20.0      22: store_ mem(v14) reg1
     1       10.0      23: reg1 = mov reg2
     1       10.0      24:  v17 = br bb4
     0        2.0      25: reg1 = load_ mem(v14)
     0        1.0      26:  v18 = ret reg1
SUM: 118.0
```

```
Final cost of the allocation:  70.0
```

From the listing above, we can easily see which instructions increased the cost significantly. These are especially `load` and `store` instructions situated in a loop.

5   Full register allocation

After the step-by-step introduction we show how to take advantage of the general register allocation procedure. It takes a function instance and a number of available registers. On success, it returns a modified copy of the input function, already after the ϕ-elimination phase. On failure, however, `None` is returned.

```
In [14]: h = bls.perform_full_register_allocation(gcd, 2)
         if h:
             print "Allocation succeeded"
             print "Cost: ", mcc.function_diff(h, gcd)

Allocation succeeded
Cost:  70.0
```

6   Comparing different algorithms

It is very easy to compare allocation algorithms. To do that, we use `utils.ResultCompSetting` class which in the constructor takes lists of functions, regcounts (different numbers of registers), allocation algorithms (allocators) and cost calculators. Then, it is passed to `utils.compute_full_results` which runs all provided algorithms on all the functions with all given numbers of registers, and calculates cost with each provided calculator. It returns a four-time list composition with all possible results, which can be printed out by `utils.compute_and_print_result_table(result, setting)`.

For example, we will compare two spilling heuristics of basic linear scan algorithm - the *Furthest-First* strategy, which spills the interval with the furthest end and *Current-First*, spilling always the current interval.

```
In [15]: import allocators.lscan.basic.spillers as spillers
         from cost import MainCostCalculator

         # Allocators
         basic_ff = BasicLinearScan(name="basic furthest first")
         basic_cf = BasicLinearScan(spiller=spillers.CurrentFirst(), name="basic current first")

         # Cost calculator
         mcc = MainCostCalculator()

         # Setup
         setting = utils.ResultCompSetting(
             inputs = m.functions.values(), # We take all the functions from the module
             regcounts = [1, 2, 3, 4, 5],
             allocators = [basic_ff, basic_cf],
             cost_calculators = [mcc])

         print "Comparing \'furthest first\' and \'current first\' strategies."
         results = utils.compute_full_results(setting)
         utils.compute_and_print_result_table(results, setting)
         print "\n"
```

Comparing 'furthest first' and 'current first' strategies.

Input	Registers	basic furthest first Main cost (S=2, N=1, L=10)	basic current first Main cost (S=2, N=1, L=10)
main	1	Failed	Failed
	2	84.0	204.0
	3	42.0	162.0
	4	0.0	0.0
	5	0.0	0.0
gcd	1	Failed	Failed
	2	70.0	151.0
	3	25.0	25.0
	4	27.0	27.0
	5	27.0	27.0

It is clear that the *Current-First* strategy is much worse than the *Furthest-First*. Small costs for higher register numbers are caused by additional `mov`-instructions inserted in the ϕ-elimination phase.

7   Regcount-to-cost plot

Another feature we may use is to plot how the cost is changing with the number of available registers. This time, we will take `SpillInstructionsCounter` instead of `MainCostCalculator` to check how fast the number of spill code decreases when we add more registers. To make the test more interesting, we will use a larger program - Fast Fourier Transform (`fft.json`). As before, we have to compute the results passing an appropriate settings object.

```python
In [16]: from cost import SpillInstructionsCounter
         sic = SpillInstructionsCounter()

         m = cfg.Module.from_file("programs/fft.json")
         analysis.perform_full_analysis(m)

         setting = utils.ResultCompSetting(
             inputs = m.functions.values(), # We take all the functions from the module
             regcounts = range(1, 25),
             allocators = [basic_ff, basic_cf],
             cost_calculators = [sic])

         results = utils.compute_full_results(setting)
         utils.plot_reg_to_cost(results, setting, figsize=(8, 8))
```

At the point (x, y) in the plot corresponding to a particular algorithm, y is the sum of costs for all functions in the module, calculated after executing the algorithm on each function with x free registers. The plot is cut for the first register numbers if there was at least one function for which the algorithm failed. Here, algorithm succeeded for all functions with just 2 registers.